



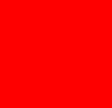
DRAFT – DO NOT CIRCULATE

ORACLE®

Learning R Series 2014

Session 2: Oracle R Advanced Analytics for Hadoop 2.3.1 – Interacting with HDFS

Mark Hornick, Director, Oracle Database Advanced Analytics



The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remain at the sole discretion of Oracle.

Topics

- What is HDFS?
- ORAAH Input Data
- ORAAH API to HDFS
 - Functions summary
 - Viewing and setting metadata
- ORAAH HDFS examples
 - Directory navigation
 - Data transfer with HDFS
 - Cleaning data and making data “pristine” data
 - Setting up the “tweets” data
- Summary

What is HDFS?

The Hadoop Distributed File System

- HDFS is the primary storage system underlying Hadoop
- Fault tolerant, scalable, highly available
- Designed to be well-suited to distributed processing
- Is superficially structured like a UNIX/Linux file system


Hadoop Distributed File System (HDFS)

- Stores data on the cluster using the native file system on *Data Nodes*
 - Loading data to HDFS is equivalent to copying files on the operating system
- Data stored as flat files
 - Automatically distributed and replicated across Data Nodes, typically 3
 - Data split into blocks, 64MB, 128MB, 256MB (default for BDA)
 - Achieves reliability and availability
- “Agreed upon” delimiters structure the data (between file and MR job)
 - Each row has (optional) key and value(s)
 - Default, tab ‘\t’ delimits key from values
 - Default, comma ‘,’ separates values
 - <line feed> indicates end of row
- MapReduce programs provide access to data on Hadoop


Distribution of Blocks across HDFS Data Nodes

Default replication factor 3

Source Data File



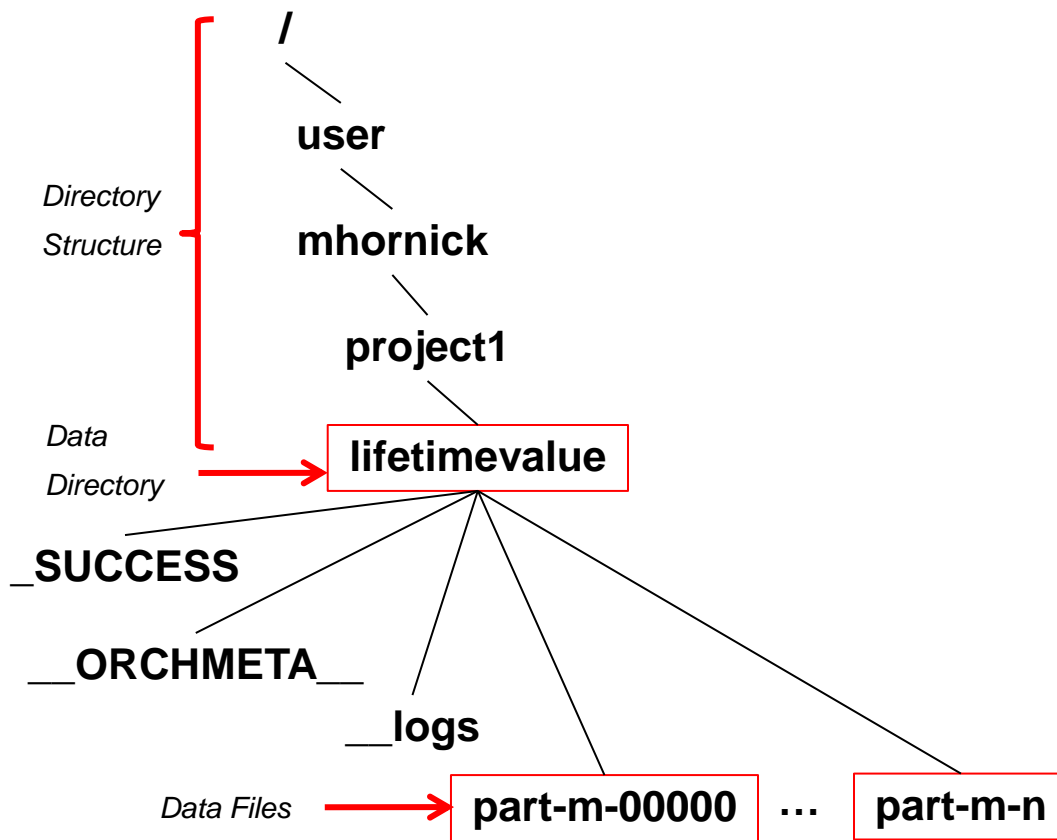
Data Nodes



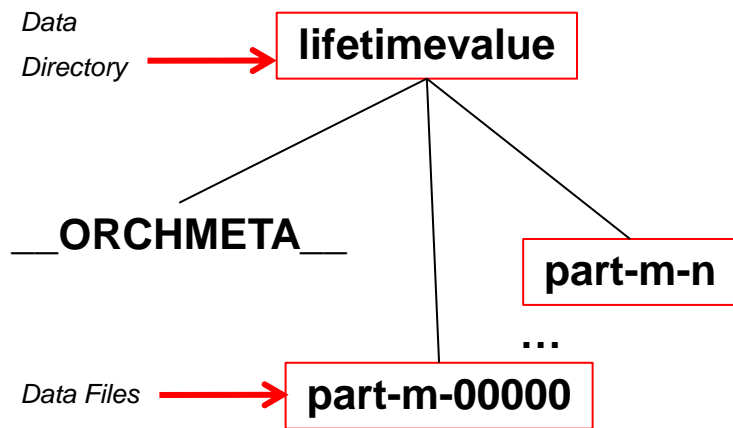
**NODE
FAILURE**

ORAAH Input Data

Structure of HDFS files



Structure of HDFS files



- Data with a key and N values
 - key<key.sep>value1<value.sep>value2...
- Data with an empty key and N values
 - <key.sep>value1<value.sep>value2...
- Data without a key and N values
 - value1<value.sep>value2...
- Data with a key and 1 value
 - key<key.sep>value
- Data without a key and 1 value
 - value
- Data with a key and no values
 - key

Data for ORAAH

- Delimited text files resident in HDFS directory or Hive tables
- `hdfs.*` functions take HDFS directories (not files) when accessing HDFS data
- ORAAH requires metadata in file `__ORCHMETA__` defined on delimited text files
 - Contains metadata about the data files
 - If `__ORCHMETA__` doesn't exist, it is created automatically during `hdfs.attach()` by sampling input files and parsing rows
 - `__ORCHMETA__` file stored alongside data files
 - With HIVE tables, `__ORCHMETA__` is auto-created from the Hive table definition

ORAAH API to HDFS

HDFS API Funtions

`hdfs.cd`

`hdfs.cleanInput`

`hdfs.cp`

`hdfs.delim`

`hdfs.describe`

`hdfs.download`

`hdfs.exists`

`hdfs.get`

`hdfs.head`

`hdfs.keysep`

`hdfs.ls`

`hdfs.meta`

`hdfs.mkdir`

`hdfs.mv`

`hdfs.ncol`

`hdfs.nrow`

`hdfs.parts`

`hdfs.pull`

`hdfs.push`

`hdfs.put`

`hdfs.pwd`

`hdfs.rm`

`hdfs.rmdir`

`hdfs.root`

`hdfs.sample`

`hdfs.setroot`

`hdfs.size`

`hdfs.sync`

`hdfs.tail`

`hdfs.upload`

`hdfs.valusep`

`is.hdfs.id`

hdfs.describe

Metadata characteristics

- **path**: Absolute HDFS path to the described object
- **origin**: description of the HDFS object origin
- **class**: R class corresponding to HDFS data
- **types**: list of data type names for each column
- **names**: vector of known column names
- **dim**: number of rows (or -1 if unknown) and columns
- **categorized**: TRUE if "factor" columns are stored as indexes
- **has.key**: TRUE if the data has key column
- **key.column**: index and name of a column containing keys
- **empty.key**: TRUE if the data has "" key
- **has.rownames**: TRUE if rownames are stored with data
- **key.sep**: delimiter used as a separator between key and values
- **value.sep**: delimiter used as a separator between values
- **quoted**: quoting symbol used when parsing fields or FALSE
- **pristine**: TRUE if data has not invalid fields
- **trimmed**: TRUE if number of columns in data can be less than "dim"

Viewing Metadata - hdfs.describe

```
> hdfs.describe("orch1cb29bfe75d")
```

	NAME	VALUE
1	path	orch1cb29bfe75d
2	origin	Uploaded "/home/mhornick/datasets/TweetsBankOfOracle.txt"
3	class	data.frame
4	types	character, logical, character, character, logical, character, numeric, character, character, character, numeric, logical, character, character
5	names	text, favorited, replyToSN, created, truncated, replyToSID, id, replyToUID, statusSource, screenName, retweetCount, retweeted, longitude, latitude
6	dim	-1 x 14
7	categorized	FALSE
8	has.key	FALSE
9	key.column	-1:NULL
10	empty.key	FALSE
11	has.rownames	FALSE
12	key.sep	\001
13	value.sep	,
14	quoted	"
15	pristine	TRUE
16	trimmed	FALSE
17	size	3745
18	parts	4

hdfs.meta

Settable characteristics

- **kvs**: Reserved of ORAAH
- **types**: Vector of type names for each column
- **names**: Vector of column names
- **class**: R class corresponding to HDFS data
- **keyi**: Index of a column containing keys
- **rownamei**: Index of a column containing row names
- **key.sep**: Symbol used as a separator between key and values
- **value.sep**: Symbol used as a separator between values
- **origin**: Description of HDFS object origin
- **dim**: Number of rows (or -1 if unknown) and columns
- **pristine**: TRUE if data has not invalid or NA values
- **quote**: Quoting symbol used for parsing data
- **categorized**: TRUE if "factor" columns are stored as indexes
- **trim**: TRUE if number of columns in data is less than "dim"

Viewing and Settings Metadata - hdfs.meta

```
> hdfs.meta("tweet_data")
$kvs
[1] TRUE
$types
[1] "character" "logical"   "character" "character" "logical"   "character" "numeric"   "character" "character"
[10] "character" "numeric"   "logical"   "character" "character"
$names
[1] "text"          "favorited"    "replyToSN"    "created"      "truncated"    "replyToSID"   "id"
[8] "replyToUID"    "statusSource" "screenName"   "retweetCount" "retweeted"     "longitude"    "latitude"
$class
[1] "data.frame"
$keyi
[1] -1
$rownamei
[1] 0
$origin
[1] "Uploaded \"/home/mhornick/datasets/TweetsBankOfOracle.txt\"
$key.sep
[1] "\\001"
$value.sep
[1] ", "
$trim
[1] FALSE
$dim
[1] -1 14
$pristine
[1] TRUE
$quote
[1] "\""
```

```
> hdfs.meta("tweet_data",pristine=TRUE)
[1] TRUE

> hdfs.ls("tweet_data")
[1] "_SUCCESS"      "__ORCHMETA__"  "_logs"
[2] "part-m-00000"  "part-m-00001"  "part-m-00002"  "part-m-00003"
```


hdfs.sync

- ORAAH maintains a cached *mini-snapshot* of HDFS metadata
- Minimizes requests to HDFS APIs to improve ORAAH HDFS function response time
- If ORAAH cache gets out of sync with current HDFS state, reset using ***hdfs.sync***
 - Deletes the metadata, forcing the re-caching of HDFS snapshot
 - If argument `dfs.id` not specified, all metadata reset
 - May need to use when an external change of HDFS object by another user or process modified HDFS content

```
x <- hdfs.put(mtcars)      # metadata is cached on write
system.time(hdfs.meta(x)) # ~0s, metadata is read from the cache
hdfs.sync(x)              # deletes cache for this object only
system.time(hdfs.meta(x)) # ~2.5s, metadata is read from HDFS
system.time(hdfs.meta(x)) # ~0s, metadata is read from the cache
```

ORAAH HDFS Examples

Demo Dataset

```
> mtcars
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4

```
...
```

```
> dim(mtcars)
```

```
[1] 32 11
```

HDFS basic functions

```
# 1. Put cars data.frame into HDFS
cars.dfs <- hdfs.put(mtcars, dfs.name="mtcars")
```

```
# 2. Create directory named 'xyz'
hdfs.mkdir("xyz")
```

```
# 3. Copy HDFS data from cars data set into xyz
HDFS directory
hdfs.cp(cars.dfs, "xyz")
```

```
# 4. List directory contents
hdfs.ls("xyz")
```

```
# 5. Remove everything under xyz
hdfs.rm("xyz/*", force=TRUE)
```

```
# 6. List directory contents
hdfs.ls("xyz")
```

```
# 7. Create directory named 'abc'
hdfs.mkdir("abc")
```

```
# 8. Move cars data into xyz
hdfs.mv(cars.dfs, "xyz")
```

```
# 9. Check existence of src HDFS directory
hdfs.exists(cars.dfs)
```

```
# 10. Move all contents of xyz into abc
hdfs.mv("xyz/*", "abc", force=TRUE)
```

```
# 11. List contents of abc
hdfs.ls("abc")
```

```
# 12. Remove directories
hdfs.rmdir("xyz")
hdfs.rmdir("abc")
```

From R to HDFS and back with categorical data

- **hdfs.put**
 - Special option "categorize", which triggers special handling of factors
 - Normally stores factors as plain strings in HDFS files
 - "categorize=TRUE" converts all factor-type columns to integers and factor level maps are written to special "sidecar" files
 - Makes representation more compact
 - Allows use of categorized data with ORAAH stats
 - Allows ORAAH to correctly restore factors when reading data back from HDFS
- **hdfs.get**
 - No new arguments
 - Automatically restore factors if stored in categorized form by hdfs.put
- **hdfs.levels**
 - Provides direct read/write access to factor levels in a map
 - Retrieve and set the mappings
 - While reading is safe, writing levels can invalidate data if not all levels specified

R → HDFS
HDFS → R

```
R> x <- hdfs.put(iris, categorize=TRUE)
R> hdfs.describe(x)
      NAME                                     VALUE
1      path                                /tmp/hdfs/tmp/orch58cd54bb77fd
2      origin                             R object "iris"
3      class                             data.frame
4      types      numeric, numeric, numeric, numeric, factor
5      names Sepal.Length, Sepal.Width, Petal.Length, Petal.Width, Species
6      dim                                150 x 5
7      categorized                        TRUE
8      has.key                           FALSE
9      key.column                         -1:NULL
10     empty.key                          FALSE
11     has.rownames                       FALSE
12     key.sep
13     value.sep                          ,
14     quoted                             FALSE
15     pristine                           TRUE
16     trimmed                           FALSE
17     size                               2558
18     parts                              1
R> hdfs.ls(x)
[1] "__ORCHLEVELS_5__" "__ORCHMETA__" "part-00000"
R> hdfs.size(x)
[1] 2558
R> y <- hdfs.put(iris)
R> hdfs.size(y)
[1] 3658
```

R → HDFS

HDFS → R

```
R> head(hdfs.get(x))
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1           5.1           3.5           1.4           0.2  setosa
2           4.9           3.0           1.4           0.2  setosa
3           4.7           3.2           1.3           0.2  setosa
4           4.6           3.1           1.5           0.2  setosa
5           5.0           3.6           1.4           0.2  setosa
6           5.4           3.9           1.7           0.4  setosa
```

```
R> hdfs.levels(x, "Species")
[1] "setosa"      "versicolor" "virginica"
R> hdfs.levels(x, Species=c("A", "B", "C"), overwrite=T)
[1] TRUE
R> hdfs.get(x)
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1           5.1           3.5           1.4           0.2      A
2           4.9           3.0           1.4           0.2      A
3           4.7           3.2           1.3           0.2      A
4           4.6           3.1           1.5           0.2      A
...
51          7.0           3.2           4.7           1.4      B
52          6.4           3.2           4.5           1.5      B
53          6.9           3.1           4.9           1.5      B
...
101          6.3           3.3           6.0           2.5      C
102          5.8           2.7           5.1           1.9      C
103          7.1           3.0           5.9           2.1      C
```

Data transfer between HDFS and file system

```
# 1. Write cars dataset to a file
write.csv(mtcars, file="mtcars.csv", row.names=FALSE)

# 2. Upload file to HDFS
cars.dfs <- hdfs.upload("mtcars.csv", header=TRUE, overwrite = TRUE, key.sep=",")

# 3. Describe HDFS object
hdfs.describe(cars.dfs)

# 4. Write the HDFS file to local disc
fileName <- hdfs.download(cars.dfs, filename="mtcars2.csv", overwrite = TRUE)

# 5. Read the local file into the R session
mtcars.new <- read.csv(fileName, header=FALSE)

# 6. Compare the two R objects
all.equal(mtcars, mtcars.new)
```

R → file system → HDFS

HDFS → file system → R

Data transfer between HDFS and database

Using scoop and Oracle Loader for Hadoop (OLH)

HDFS → DB
DB → HDFS

```
hdfs.toDB <- hdfs.pull  
hdfs.fromDB <- hdfs.push
```

```
# 1. Copy data from HDFS to DB using sqoop (default)  
hdfs.toDB(dfs.id = cars.dfs,  
          db.name = "MTCARS", overwrite = TRUE)
```

```
# 2. Check class of returned ORACLE DB object  
ore.sync(table="MTCARS")  
ore.attach()  
class(MTCARS)
```

```
# 3. Push data from DB into HDFS using sqoop  
cars.dfs.id1 <- hdfs.fromDB( MTCARS,  
                            dfs.name = "cars_dfs1",  
                            overwrite = TRUE)
```

```
hdfs.describe(cars.dfs.id1)  
hdfs.get(cars.dfs.id1)
```

```
# 4. Copy data from HDFS data to DB using OLH  
# Requires OLH.  
hdfs.toDB(dfs.id = cars.dfs,  
          db.name = "MTCARS2",  
          overwrite = TRUE, driver = "olh")
```

```
# 5. Check class of returned DB object  
ore.sync(table="MTCARS2")  
ore.attach()  
class(MTCARS2)
```

```
# 6. Drop database tables  
ore.drop(table=c("MTCARS", "MTCARS2"))
```

HDFS directory navigation

1. Check present working directory

```
hdfs.pwd()
```

2. Create HDFS directory named xyz

```
hdfs.mkdir("xyz")
```

3. Change directory to xyz

```
hdfs.cd("xyz")
```

4. Create HDFS dir 'abc' in 'xyz'

```
hdfs.mkdir("abc")
```

5. Change directory to abc

```
hdfs.cd("abc")
```

6. List current directory

```
hdfs.ls()
```

7. Go to parent directory

```
hdfs.cd('..')
```

8. List contents of current directory

```
hdfs.ls()
```

9. Go to parent directory

```
hdfs.cd('..')
```

10. cd using absolute path

```
hdfs.cd(file.path(hdfs.pwd(), '/xyz/abc'))
```

11. cd to HDFS directory 'xyz'

```
hdfs.cd('..')
```

12. try removing all contents of xyz

```
hdfs.rmdir('*')
```

13. List contents of current directory

```
hdfs.ls()
```

14. Use force to do above w/no prompt

```
hdfs.rmdir('*', force=TRUE)
```

15. Go to parent directory

```
hdfs.cd('..')
```

16. Get present working directory

```
my.pwd <- hdfs.pwd()
```

17. cd to hdfs root ('/')

```
hdfs.cd()
```

18. Remove xyz w/absolute path

```
hdfs.rmdir(file.path(my.pwd, "xyz"))
```

19. Restore working directory

```
hdfs.cd(my.pwd)
```

hdfs.cleanInput

- Makes data “pristine” and sets metadata pristine=TRUE
- Remove invalid values or replace them with default values
- `hdfs.cleanInput(input, config = NULL, tmpdir = "/tmp",
replace = TRUE, replace_val = NULL)`
- Returns ORAAH HDFS identifier of cleaned output
- Displays
 - Number of cells replaced when replace = TRUE
 - Number of rows removed when replace = FALSE
 - Percentage of cells replaced when replace = TRUE
 - Percentage of rows removed when replace = FALSE
 - Total number of input rows

Cleaning HDFS data using ORAAH with defaults

```
# Create data.frame with some missing values
df <- data.frame(x=c(1,2,3,4,5,6), y=c(1,2,3,NA,NA,6))

df.dfs <- hdfs.put(df)                                # Write data.frame to HDFS

df.dfs.clean <- hdfs.cleanInput(df.dfs) # substitute NAs with 0s (default value)

hdfs.get(df.dfs.clean)                                # Get transformed output
hdfs.rm(df.dfs.clean)

# Clean input by removing rows with missing values (NAs)
df.dfs.clean <- hdfs.cleanInput(df.dfs, replace = FALSE)

hdfs.get(df.dfs.clean)
hdfs.rm(df.dfs.clean)
hdfs.rm(df.dfs)
```

Cleaning HDFS data using ORAAH with custom values

```
# Create data.frame with numeric and character columns containing missing values
df <- data.frame(x=c(1,NA,NA,4,5,6),
                 y=c("abc","def","efg",NA,NA,"xyz"), stringsAsFactors=FALSE)

df.dfs <- hdfs.put(df)

# Substitute numeric NAs with -1 and character NAs with "abc"
df.dfs.clean <-
  hdfs.cleanInput(df.dfs, replace_val = data.frame(numeric=-1, character="abc",
                                                    stringsAsFactors=FALSE))



hdfs.get(df.dfs.clean)      # Retrieve and print transformed output

hdfs.rm(df.dfs.clean)
hdfs.rm(df.dfs)
```

Pristine mode

- Provides major performance enhancement
 - Avoids validating data on read
- "Pristine" metadata attribute defines data as
 - Every row having same number of columns
 - Missing values represented either as "NA" or empty string ""
 - No non-numeric values in numeric columns, except for missing values
- `hdfs.meta(x, pristine=TRUE)`

Performance benefits of Pristine Mode



Using ORAAH Pristine Mode

Demonstrate performance improvement using pristine data

Create a 100 MB dataset with 20% NA values

```
data_NA20 <- orch.datagen(datasize=1.2e+8, map.degree=5, numeric.col.count = 200,  
                          percent.na=20)
```

```
hdfs.describe(data_NA20)
```

Create output metadata string for the hadoop job

```
meta_str <- sprintf("data.frame(%s)", paste0("val",1:200,"=0", collapse=","))
```

Time simple mapper-only job performing sampling in R

```
system.time(x <- hadoop.run(  
  data_NA20,  
  mapper = function(key, val) {  
    select <- (runif(nrow(val)) <= (percent/100))  
    orch.keyvals(key[select], val[select,])  
  },  
  export = orch.export(percent=1),  
  config = new("mapred.config",  
    map.output = eval(parse(text=meta_str))) # mapper output metadata  
))
```


Using ORAAH Pristine Mode

Demonstrate performance improvement using pristine data

```
hdfs.rm(x)

# Assign metadata to indicate data is pristine
data_NA20p <- hdfs.meta(data_NA20, pristine=TRUE)

# Time simple mapper-only job on pristine input
system.time(x <- hadoop.run(
  data_NA20p,
  mapper = function(key, val) {
    select <- runif(nrow(val)) <= (percent/100)
    orch.keyvals(key[select], val[select,])
  },
  export = orch.export(percent=1),
  config = new("mapred.config",
    map.output = eval(parse(text=meta_str)))
))
```

Sampling HDFS and Hive data

Using the built-in `orch.sample` function

```
cars.dfs <- hdfs.put(mtcars, dfs.name="/tmp/cars_tmp")
cars.dfs.samp <- orch.sample(cars.dfs, percent = 10, output="/tmp/samp_out10")
hdfs.get(cars.dfs.samp)
```

```
cars.dfs.samp.r <- orch.sample(cars.dfs, nrows = 20, output="/tmp/samp_out20r")
hdfs.get(cars.dfs.samp.r)
```

```
ore.create(mtcars, table="cars1hive")    # Create HIVE table cars1hive from mtcars
cars.hive.samp1 <- orch.sample(cars1hive, percent = 10)
cars.hive.samp1
```

```
cars.hive.samp2 <- orch.sample(cars1hive, percent = 10, output="samp_out10")
cars.hive.samp2
```

```
cars.hive.samp2.r <- orch.sample(cars1hive, nrows = 20, output="samp_out20r")
cars.hive.samp2.r
```

Tweets

```
"text","favorited","replyToSN","created","truncated","replyToSID","id","replyToUID",  
  "statusSource","screenName","retweetCount","retweeted","longitude","latitude"  
"Doing a great job #SavingsAlpha #BankOfOracle #SavingsBeta",FALSE,NA,2014-01-01  
  00:00:00,FALSE,NA,3.430311e+17,NA,"<a href='\"http://www.hootsuite.com\"'  
  rel='\"nofollow\"'>HootSuite</a>","MEE.COMER.CU1142",0,FALSE,NA,NA  
"Where can I get #SavingsBeta #BankOfOracle",FALSE,NA,2014-01-01  
  03:40:28,FALSE,NA,3.430311e+17,NA,"<a href='\"http://accounts.vitrue.com/\"'  
  rel='\"nofollow\"'>Vitrue Accounts</a>","LAURINDA.ROWLAND.CU1144",0,FALSE,NA,NA  
"I'm a fan of #BOOCD #SavingsBeta #SavingsAlpha",FALSE,NA,2014-01-01  
  07:20:57,FALSE,NA,3.430311e+17,NA,"web","ANNETT.MCMULLEN.CU1145",0,FALSE,NA,NA  
"I'm a fan of #BankOfOracle #SavingsBeta #SavingsAlpha",FALSE,NA,2014-01-01  
  11:01:26,FALSE,NA,3.430311e+17,NA,"<a href='\"http://www.tweetcaster.com\"'  
  rel='\"nofollow\"'>TweetCaster for Android</a>","THELMA.DELONG.CU1146",0,FALSE,NA,NA  
"Where can I get #CheckingPlusPlus",FALSE,NA,2014-01-01  
  14:41:55,FALSE,NA,3.430311e+17,NA,"<a href='\"http://www.tweetdeck.com\"'  
  rel='\"nofollow\"'>TweetDeck</a>","CRISELDA.HAWKINS.CU1147",1,FALSE,NA,NA  
"
```

Tweet Example – Loading and setting up data

```
tweets.id <- hdfs.upload("~/datasets/tweets.txt",dfs.id="tweets",
                        header=FALSE,overwrite=TRUE,key.sep='\1',value.sep=',') # bogus key.sep = no key

hdfs.meta(tweets.id, names=c("text","favorited","replyToSN","created",
                             "truncated","replyToSID","id","replyToUID","statusSource","screenName",
                             "retweetCount","retweeted","longitude","latitude"))
hdfs.meta(tweets.id, pristine=TRUE, quote='')
hdfs.meta(tweets.id)

tweets.1000 <- orch.sample(tweets.id, percent=1,output="tweetsBOO.1000")
tweets.20   <- orch.sample(tweets.1000,percent=2,output="tweetsBOO.20")
```

Summary

- ORAAH enables creation, manipulation, and viewing of HDFS data
- Specialized functions enable import/export of data
 - HDFS \leftrightarrow R
 - HDFS \leftrightarrow Database
 - HDFS \leftrightarrow File System
 - HDFS \leftrightarrow Hive
- Supports automatic discovery of metadata
- Performance optimized via caching of metadata

Resources

<http://www.oracle.com/goto/R>

- **Blog:** <https://blogs.oracle.com/R/>
- **Forum:** <https://forums.oracle.com/forums/forum.jspa?forumID=1397>
- **Book:** [Using R to Unlock the Value of Big Data](#) Oracle Press
- **Oracle R Distribution:**
<http://www.oracle.com/technetwork/indexes/downloads/r-distribution-1532464.html>
- **ROracle:**
<http://cran.r-project.org/web/packages/ROracle>
- **Oracle R Enterprise:**
<http://www.oracle.com/technetwork/database/options/advanced-analytics/r-enterprise>
- **Oracle R Advanced Analytics for Hadoop:**
<http://www.oracle.com/us/products/database/big-data-connectors/overview>



ORACLE®