



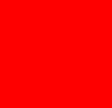
DRAFT – DO NOT CIRCULATE

ORACLE®

Learning R Series 2014

Session 3: Oracle R Advanced Analytics for Hadoop 2.3.1 – Using ORCHhive

Mark Hornick, Director, Oracle Database Advanced Analytics



The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remain at the sole discretion of Oracle.

Topics

- What is Hive?
- ORCHhive Transparency Layer
- Examples
- Summary

What is Hive?

What is Hive?

- SQL-like abstraction on Hadoop
- Becoming *de facto* standard for SQL based apps on Hadoop
- Converts SQL queries to MapReduce jobs to be run on Hadoop
- Provides simple query language (HQL) based on SQL
- Enables non-Java users to leverage Hadoop via SQL-like interfaces



Motivation for ORCHhive

- “Big data” scalability and performance for R users on Hadoop
 - Enable R users to clean, explore, and prepare HIVE data transparently
 - Ready data for analytic techniques using ORAAH MapReduce framework
-
- ORE provides transparent access to database tables and views from R based on SQL mapping
 - Since Hive is SQL-based, it is a natural extension to provide ORE-type transparency on top of Hive HQL for R users to HDFS data

When to use ORCHhive

- Need to prepare Hive data
 - Row filtering
 - Column projection
 - Missing value handling
- Analyze Hive data
 - Mathematical functions
 - Aggregation functions
- Desire to use standard R data.frame-type functions
 - Corresponds to ORE transparency layer
- Data in Hive (or HDFS and can be converted to Hive)

ORCHhive Transparency Layer

Working with ORAAH Hive interface

- Available as *special* data frames – ore.frame
- Objects contain enough metadata to generate HiveQL queries when data in Hive tables/views are processed
- Supports Hive tables/views from default and non-default Hive databases

Supported R Functions

- **Storage methods**
 - ore.create, ore.drop, ore.push, ore.pull, ore.get
- **Methods**
 - is.ore.frame, is.ore.vector, is.ore.logical, is.ore.integer, is.ore.numeric, is.ore.character, is.ore, as.ore.frame, as.ore.vector, as.ore.logical, as.ore.integer, as.ore.numeric, as.ore.character, as.ore
- **ore.vector methods**
 - show, length, c, is.vector, as.vector, as.character, as.numeric, as.integer, as.logical, "[", "[<-", l, Compare, ore.recode, is.na, "%in%", unique, sort, table, paste, tapply, by, head, tail
- **ore.logical methods**
 - <, >, ==, <=, >=, !, xor, ifelse, and, or
- **ore.number methods**
 - +, -, *, ^, %%, %/%, /, is.finite, is.infinite, is.nan, abs, sign, sqrt, ceiling, floor, trunc, log, log10, log2, log1p, logb, acos, asin, atan, exp, expm1, cos, sin, tan, zapsmall, round, Summary, summary, mean
- **ore.character methods**
 - nchar, tolower, toupper, casefold, gsub, substr, substring
- **ore.frame methods**
 - show, attach, [, \$, \$<-, [[, [[<-, head, tail, length, nrow, ncol, NROW, NCOL, dim, names, names<-, colnames, colnames<-, as.list, unlist, summary, rbind, cbind, data.frame, as.data.frame, as.env, eval, +, -, *, ^, %%, %/%, /, Compare, Logic, !, xor, is.na, is.finite, is.infinite, is.nan, abs, sign, sqrt, ceiling, floor, trunc, log, log10, log2, log1p, logb, acos, asin, atan, exp, expm1, cos, sin, tan, round, Summary, rowSums, colSums, rowMeans, colMeans, unique, by, merge
- **Aggregate functions**
 - OREStats: fivenum, aggregate, quantile, sd, var(only for vectors), median, IQR

Hive tips

- Hive table and column names will be converted to lower case
- Table and column names cannot contain “.” – will be converted to “_”
- Does not support factors – must convert to character vector first to create table
- Beware of Hive keywords, like “FIRST” as column names, these must be renamed
- Row.names are not supported in ORCHhive, so functions requiring ordering not available

ORCHhive Introductory Example

Example using OREhive

```
ore.connect(type="HIVE")
```

```
ore.attach()
```

```
# create a Hive table by pushing the numeric
```

```
# columns of the iris data set
```

```
iris_table <- ore.push(iris[1:4])
```

```
# Create bins based on Petal Length
```

```
iris_table$PetalBins =
```

```
  ifelse(iris_table$Petal.Length < 2.0, "SMALL PETALS",
```

```
+ ifelse(iris_table$Petal.Length < 4.0, "MEDIUM PETALS",
```

```
+ ifelse(iris_table$Petal.Length < 6.0,
```

```
+ "MEDIUM LARGE PETALS", "LARGE PETALS")))
```

```
#PetalBins is now a derived column of the HIVE object
```

```
> names(iris_table)
```

```
[1] "Sepal.Length" "Sepal.Width" "Petal.Length"
```

```
[4] "Petal.Width" "PetalBins"
```

```
# Based on the bins, generate summary statistics for each group
```

```
aggregate(iris_table$Petal.Length,
```

```
  by = list(PetalBins = iris_table$PetalBins),
```

```
+ FUN = summary)
```

```
1      LARGE PETALS      6 6.025000 6.200000 6.354545 6.612500 6.9 0
```

```
2 MEDIUM LARGE PETALS  4 4.418750 4.820000 4.888462 5.275000 5.9 0
```

```
3      MEDIUM PETALS    3 3.262500 3.550000 3.581818 3.808333 3.9 0
```

```
4      SMALL PETALS     1 1.311538 1.407692 1.462000 1.507143 1.9 0
```

```
Warning message:
```

```
ORE object has no unique key - using random order
```

ORCHhive Interface

ORCHhive basic functions

```
ore.connect(type="HIVE") # Connect to HIVE
ore.attach()             # Attach the current envt. into search path of R
iris_temp <- iris
summary(iris_temp)
colMeans(iris_temp[,1:4])

# ORCHhive does not support factors yet, convert factor columns to character
factfilt <- sapply(iris_temp, is.factor)
iris_temp[factfilt] <- data.frame(lapply(iris_temp[factfilt], as.character),
                                stringsAsFactors = FALSE)

iris_hive_table <- ore.push(iris_temp) # Push iris_temp data.frame to temp Hive table
class(iris_hive_table)
summary(iris_hive_table)
colMeans(iris_hive_table[,1:4])

ore.create(iris_hive_temp, table="iris_table") # create persistent Hive table
```

ORCHhive basics

```
# Number of rows
nrow(iris_table)

# Column names of the data frame
names(iris_table)

# Number of columns of the data frame
length(iris_table)
ncol(iris_table)

# Head
head(iris_table, n = 5)
```

```
# Vectors
class(iris_table$Sepal.Length)

class(iris_table$Species)

# 'is' variants
is.character(iris_table$Species)
is.ore.character(iris_table$Species)

# Number of characters in each column value
nchar(iris_table$Species)
```


ORCHhive Aggregate

```
for(fun in c("length","summary","mean","min","max ","sd","median","IQR")) {  
  x = aggregate(iris_table$Petal.Length, by = list(Species = iris_table$Species),  
                FUN = fun)  
  class(x)  
  x  
}  
  
# More than one grouping column  
x = aggregate(iris_table$Petal.Length,  
              by = list(Species = iris_table$Species,  
                        width = iris_table$Petal.Width),  
              FUN = length)  
x
```

ORCHhive analysis

```
# What are the unique Species?
```

```
unique(iris_table$Species)
```

```
# Number of unique Species
```

```
length(unique(iris_table$Species))
```

```
# Count of observations with Species = "setosa"
```

```
nrow(iris_table[iris_table$Species == "setosa", ])
```

```
# Count of rows where Species == "setosa" and Petal.Width == 0.3
```

```
nrow(iris_table[iris_table$Species == "setosa" & iris_table$Petal.Width == 0.3, ])
```

ORCHhive row filtering

```
# On an ore.frame the result is just a logical query
iris_temp_new = iris_table[iris_table$Petal.Width <= 0.3, ]
nrow(iris_temp_new)
class(iris_temp_new)

# Missing is NA in R, count observations where Petal.Length is missing
nrow(iris_table[is.na(iris_table$Petal.Length), ])

# Or the other way round...
nrow(iris_table[!is.na(iris_table$Petal.Length), ])
```

ORCHhive binning

```
# Create bins based on Petal Length
iris_table$PetalBins =
  ifelse(iris_table$Petal.Length < 2.0, "SMALL PETALS",
    ifelse(iris_table$Petal.Length < 4.0, "MEDIUM PETALS",
      ifelse(iris_table$Petal.Length < 6.0, "MEDIUM LARGE PETALS", "LARGE PETALS")))
iris_table

# Get frequency count per species
table(iris_table$Species, iris_table$PetalBins)

# Reusable binning logic - a.k.a. FORMATS in SAS
PetalBins = function(x) {ifelse(x < 2.0, "SMALL PETALS",
  ifelse(x < 4.0, "MEDIUM PETALS",
    ifelse(x < 6.0, "MEDIUM LARGE PETALS", "LARGE PETALS")))}

PetalBins(iris_table$Petal.Length)
```

ORCHhive using non-default database

```
ore.drop(table="iris_table")
# create a non-default HIVE database
try(ore.exec("drop database dbtmp"),
    silent = TRUE)
ore.exec("create database dbtmp")

# point R environment to this
ore.hiveOptions(dbname='dbtmp')
ore.showHiveOptions()

# create a table in this HIVE database
ore.create(iris[1:4], table="iris_table")

# restore HIVE environment to default
ore.hiveOptions()
ore.showHiveOptions()
```

```
# error since iris_table not in default db
errmsg <- try(nrow(iris_table),
              silent = TRUE)

# show error
errmsg

# move back to dbtmp database
ore.hiveOptions(dbname='dbtmp')
# this should succeed
nrow(iris_table)

# cleanups
ore.drop(table="iris_table")
# restore HIVE defaults
ore.hiveOptions()
# drop HIVE database
ore.exec("drop database dbtmp")
```

ORCHhive column functions

```
# Select distinct(Species)
unique(iris_table$Species)

x <- iris_table$Petal.Length
min(x)
max(x)
sd(x)
mean(x)
fivenum(x)
var(x)
IQR(x)
quantile(x)

log(x)
log10(x)
log2(x)
```

```
abs(x)
sqrt(x)

exp(x)
expm1(x)
round(x)

# String Functions
substr(iris_table$Species, 1, 2)
tolower(iris_table$Species)
toupper(iris_table$Species)
x = gsub("s", "v", iris_table$Species)
x
```

ORCHhive NULLs

```
AIRQUALITY <- ore.push(airquality)      # Push airquality data.frame to HIVE
class(AIRQUALITY)

# Explore NA behavior of R
nrow(airquality_temp[airquality_temp$Ozone < 30,])

# Explicit exclusion of NAs
nrow(airquality_temp[airquality_temp$Ozone < 30 & !is.na(airquality_temp$Ozone),])

# Default HIVE table behavior: exclude NULLS in output
nrow(AIRQUALITY[AIRQUALITY$Ozone < 30,])

# For R's NA behavior, request it explicitly
options(ore.na.extract = TRUE)
nrow(AIRQUALITY[AIRQUALITY$Ozone < 30,])
```

HIVE NULLS - results

```
> AIRQUALITY <- ore.push(airquality) # Push airquality data.frame to HIVE
> class(AIRQUALITY)
[1] "ore.frame"
attr(,"package")
[1] "OREbase"
>
> # Explore NA behavior of R
> nrow(airquality_temp[airquality_temp$Ozone < 30,]) # Return all observations where ozone < 30
[1] 92
>
> # Explicit exclusion of NAs
> nrow(airquality_temp[airquality_temp$Ozone < 30 & !is.na(airquality_temp$Ozone),])
[1] 55
>
> nrow(AIRQUALITY[AIRQUALITY$Ozone < 30,]) # Default HIVE table behavior: exclude NULLS in output
[1] 55
>
> options(ore.na.extract = TRUE) # For R's NA behavior, request it explicitly
> nrow(AIRQUALITY[AIRQUALITY$Ozone < 30,])
[1] 92
```


Push and Pull

- ore.push creates temporary ore.frame in HIVE table
- ore.pull materializes HIVE table in R client memory as data.frame
- Factors not yet supported, convert to character type

HIVE push/pull

```
iris_temp <- iris                                # create copy of iris dataset

factfilt <- sapply(iris_temp, is.factor)          # Convert factor cols to character
iris_temp[factfilt] <- data.frame(lapply(iris_temp[factfilt], as.character),
                                   stringsAsFactors = FALSE)

iris_table <- ore.push(iris_temp)                 # Push iris_temp data frame to HIVE
class(iris_table)

iris_table2 <- iris_table[iris_table$Species == "setosa", ] # Filter one Species
inmem_df    <- ore.pull(iris_table2)              # Pull filtered rows to R client
class(inmem_df)

(r_model <- lm(Petal.Length ~ Petal.Width, inmem_df)) # Apply R lm on inmem_df
df       <- data.frame(residuals = residuals(r_model))

(RESIDUALS <- ore.push(df))                       # Push df to HIVE
```

HIVE push/pull - results

```
> iris_table <- ore.push(iris_temp)                                # Push iris_temp data frame to HIVE
> class(iris_table)
[1] "ore.frame"
attr(,"package")
[1] "OREbase"
>
> iris_table2 <- iris_table[iris_table$Species == "setosa", ] # Filter one Species
> inmem_df <- ore.pull(iris_table2)                                # Pull filtered rows to R client
Warning message:
ORE object has no unique key - using random order
> class(inmem_df)
[1] "data.frame"
>
> (r_model <- lm(Petal.Length ~ Petal.Width, inmem_df))           # Apply R lm on inmem_df
Call:
lm(formula = Petal.Length ~ Petal.Width, data = inmem_df)
Coefficients:
(Intercept)  Petal.Width
      1.3276       0.5465

> df <- data.frame(residuals = residuals(r_model))
> (RESIDUALS <- ore.push(df))                                     # Push df to HIVE
      residuals
1 -0.036861448
2 -0.036861448 ...
```

HIVE Sequence Files

- Use STORED AS SEQUENCEFILE if the data needs to be compressed
- Benefits
 - Reduce disk space usage
 - Can yield better query performance than uncompressed storage
- A SequenceFile can be split by Hadoop and distributed across map jobs

HIVE Sequence Files

```
ore.drop(table=c("cars_tab", "cars_seq")) # Remove HIVE tables if they exist

ore.create(cars, table="cars_tab")          # Create a HIVE table from R's cars dataset

cars$speed                                # Print 'speed' of the cars data.frame
cars_tab$speed                             # Print 'speed' for the cars_tab HIVE table

# Execute a HIVE DDL to create a HIVE table from cars_tab but stored as sequencefile
ore.exec("create table cars_seq stored as sequencefile as select * from cars_tab")

ore.sync(table="cars_seq") # Sync the table created above

cars_seq$speed                             # Print the column named 'speed' for cars_seq HIVE table

cars_seq$speed + 1                          # Add one to above column
```

HIVE to HDFS and back

```
hdfs.fromHIVE <- orch.hive2hdfs
hdfs.toHIVE <- orch.hdfs2hive

ore.drop(table="CEMENT_TAB")           # Remove HIVE table if it exists
ore.create(cement, table="CEMENT_TAB")  # Create a HIVE table from R data set cement

filtered_x <- CEMENT_TAB[CEMENT_TAB$x1 > 7 & CEMENT_TAB$x4 < 45, ] # Filter rows, project columns
filtered_x <- filtered_x[, c('x1', 'x2', 'x3')]

dfs.id <- hdfs.fromHIVE(filtered_x)      # Convert transformed ore.frame to HDFS entity
hdfs.get(dfs.id)                         # Use ORAAH functions to view data
hdfs.size(dfs.id)

filtered_x_hive <- hdfs.toHIVE(dfs.id)    # Convert HDFS data to HIVE table

nrow(filtered_x_hive)                   # use transparency layer functions on HIVE table
head(filtered_x_hive)

hdfs.rm(dfs.id)                          # Cleanup
ore.drop(table="CEMENT_TAB")
```

ORCHhive Summary

- Convenient way to manipulate data in Hadoop from R
- Avoids writing MapReduce jobs to process data
- Avoids learning HiveQL to process data
- Since Hive is SQL-based, ORCHhive is a natural extension to provide ORE-type transparency on top of Hive HQL for R users to HDFS data

Resources

<http://www.oracle.com/goto/R>

- **Blog:** <https://blogs.oracle.com/R/>
- **Forum:** <https://forums.oracle.com/forums/forum.jspa?forumID=1397>
- **Book:** [Using R to Unlock the Value of Big Data](#) Oracle Press
- **Oracle R Distribution:**
<http://www.oracle.com/technetwork/indexes/downloads/r-distribution-1532464.html>
- **ROracle:**
<http://cran.r-project.org/web/packages/ROracle>
- **Oracle R Enterprise:**
<http://www.oracle.com/technetwork/database/options/advanced-analytics/r-enterprise>
- **Oracle R Advanced Analytics for Hadoop:**
<http://www.oracle.com/us/products/database/big-data-connectors/overview>



ORACLE®